

Delhi Technological University

(Formerly Delhi College of Engineering)

Shahbad Daulatpur, Bawana Road, Delhi-110042

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING



Artificial Neural Networks

Subject Code: AI504

Lab File

Submitted To:

Anil Singh Parihar Sir

Department of Computer Science

And Engineering

Submitted By:

SHIKHAR ASTHANA

M.Tech AFI

2K22/AFI/24

INDEX

S. No.	Objective	Date	Sign
1	To study about NumPy, Pandas, and Matplotlib libraries in python.		
2	To perform data preprocessing and data visualization on iris dataset.		
3	To implement uni-variate linear regression using gradient descent.		
4	To implement multi-variate linear regression using gradient descent.		
5	To implement logistic regression using gradient descent.		
6	To implement basic perceptron model.		
7	To implement multi-layer neural network.		
8	To perform classification on MNIST Digit dataset using Neural Network.		
9	To perform classification on MNIST Fashion dataset using Neural Network.		

EXPERIMENT 1

AIM: To study about NumPy, Pandas, and Matplotlib libraries in python.

THEORY: NumPy is a library for the Python programming language, adding support for large, multi-dimensional arrays and matrices, along with a large collection of high-level mathematical functions to operate on these arrays. NumPy is widely used for scientific computing and data analysis tasks.

Pandas is a Python library for data manipulation and analysis. It provides data structures for efficiently storing and manipulating large datasets and tools for working with structured data, such as SQL tables or Excel spreadsheets.

Matplotlib is a plotting library for the Python programming language and its numerical mathematics extension NumPy. It provides an object-oriented API for embedding plots into applications using general-purpose GUI toolkits like Tkinter, wxPython, Qt, or GTK.

CODE & OUTPUT:

The screenshot shows a Jupyter Notebook interface with the following details:

- Header:** jupyter Lab1_NumpyPandasMatplotlib Last Checkpoint: a few seconds ago (autosaved) | Logout
- Toolbar:** File, Edit, View, Insert, Cell, Kernel, Help, Not Trusted, Python 3 C
- In [1]:**

```
1 #####  
2 # LAB 1 - NumPy, Pandas, and Matplotlib Libraries #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####
```
- In [2]:**

```
1 #importing the Libraries  
2 import numpy as np  
3 import pandas as pd  
4 import matplotlib.pyplot as plt
```
- In [3]:**

```
1 #Exploring NumPy Library  
2  
3 a = np.array([1,2,3,4,5])  
4 b = np.array([[1,2,3],[4,5,6]])  
5  
6 print(f'a dimensions: {a.ndim}')  
7 print(f'a shape: {a.shape}')  
8 print(f'a length: {len(a)}')  
9  
10 print(f'b dimensions: {b.ndim}')  
11 print(f'b shape: {b.shape}')  
12 print(f'b length: {len(b)})'
```

Output:
a dimensions: 1
a shape: (5,)
a length: 5
b dimensions: 2
b shape: (2, 3)
b length: 2
- In [4]:**

```
1 c = np.ones((3,3))  
2 d = np.zeros((3,5))  
3 print(c)  
4 print(d)
```

Output:
[[1. 1. 1.]
 [1. 1. 1.]
 [1. 1. 1.]]
[[0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]]

```
In [5]:
```

```
1 e = np.eye(3)
2 print(e)
```

```
[[1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 1.]]
```

```
In [6]:
```

```
1 #operations in numpy
2 f = np.dot(c,d)
3 print(f)
```

```
[[0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0.]]
```

```
In [7]:
```

```
1 print(np.logical_or(np.array([1,0,0,0]),np.array([1,1,0,0])))
2 print(np.logical_and(np.array([1,0,0,0]),np.array([1,1,0,0])))
```

```
[ True  True False False]
[ True False False False]
```

```
In [8]:
```

```
1 #Exploring Pandas Library
2
3 #importing the iris dataset CSV file using pandas
4 iris_df = pd.read_csv("iris.csv")
```

```
In [9]:
```

```
1 #Exploratory commands of dataframe in pandas
2
3 iris_df.head()
```

```
Out[9]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
In [10]:
```

```
1 iris_df.tail()
```

```
Out[10]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
145	146	6.7	3.0	5.2	2.3	Iris-virginica
146	147	6.3	2.5	5.0	1.9	Iris-virginica
147	148	6.5	3.0	5.2	2.0	Iris-virginica
148	149	6.2	3.4	5.4	2.3	Iris-virginica
149	150	5.9	3.0	5.1	1.8	Iris-virginica

```
In [11]:
```

```
1 iris_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   Id          150 non-null    int64  
 1   SepalLengthCm 150 non-null  float64 
 2   SepalWidthCm  150 non-null  float64 
 3   PetalLengthCm 150 non-null  float64 
 4   PetalWidthCm  150 non-null  float64 
 5   Species      150 non-null   object  
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

```
In [12]:
```

```
1 iris_df.columns
```

```
Out[12]: Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',
 'Species'],
 dtype='object')
```

```
In [13]:
```

```
1 iris_df["Species"].unique()
```

```
Out[13]: array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

```
In [14]: 1 #Filtering commands of dataframes in pandas  
2  
3 iris_df[iris_df["Species"]=="Iris-setosa"].head()
```

Out[14]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

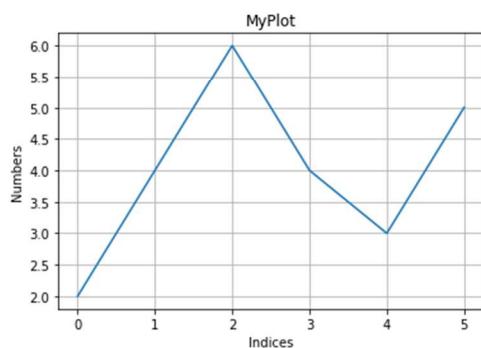
```
In [15]: 1 iris_df[(iris_df["Species"]=="Iris-setosa") | (iris_df["Species"]=="Iris-versicolor")]
```

Out[15]:

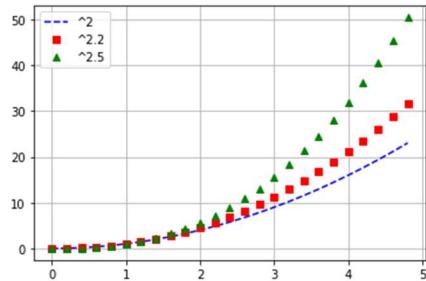
	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...
95	96	5.7	3.0	4.2	1.2	Iris-versicolor
96	97	5.7	2.9	4.2	1.3	Iris-versicolor
97	98	6.2	2.9	4.3	1.3	Iris-versicolor
98	99	5.1	2.5	3.0	1.1	Iris-versicolor
99	100	5.7	2.8	4.1	1.3	Iris-versicolor

100 rows × 6 columns

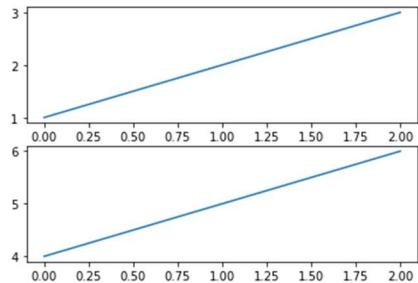
```
In [16]: 1 #Exploring Matplotlib  
2  
3 plt.plot([2,4,6,4,3,5])  
4 plt.ylabel("Numbers")  
5 plt.xlabel("Indices")  
6 plt.title("MyPlot")  
7 plt.grid()  
8 plt.show()
```



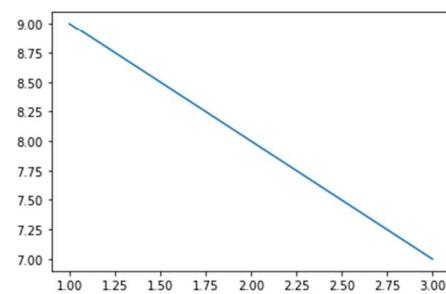
```
In [17]: 1 t = np.arange(0,5,0.2)
2
3 plt.plot(t,t**2,'b--',label="^2")
4 plt.plot(t,t**2.2,'rs',label="^2.2")
5 plt.plot(t,t**2.5,'g^',label="^2.5")
6 plt.grid()
7 plt.legend()
8 plt.show()
```



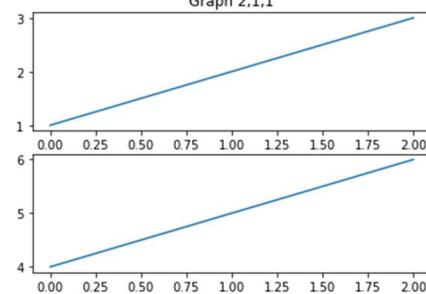
```
In [18]: 1 plt.figure(1)
2 plt.subplot(2,1,1)
3 plt.plot([1,2,3])
4 plt.subplot(2,1,2)
5 plt.plot([4,5,6])
6 plt.show()
```



```
In [19]: 1 plt.figure(1)
2 plt.subplot(2,1,1)
3 plt.plot([1,2,3])
4 plt.subplot(2,1,2)
5 plt.plot([4,5,6])
6
7 plt.figure(2)
8 plt.plot([1,2,3],[9,8,7])
9 plt.figure(1)
10 plt.subplot(2,1,1)
11 plt.title("Graph 2,1,1")
12 plt.show()
```



Graph 2,1,1



LEARNING OUTCOMES:

- Gain a better understanding of the NumPy, Pandas, and Matplotlib libraries.
- Learn how to use these libraries for data manipulation, analysis, and visualization.
- Understand how to use NumPy arrays, Pandas data frames, and Matplotlib plots to represent and manipulate data in Python.

EXPERIMENT 2

AIM: To perform data preprocessing and data visualization on iris dataset.

THEORY:

- Data preprocessing is the process of cleaning, transforming, and preparing raw data for analysis. It involves techniques such as data cleaning, feature scaling, feature extraction, and data integration.
- Data visualization is the representation of data in a graphical or pictorial format. It is used to help understand and communicate patterns, trends, and relationships in data.
- The iris dataset is a classic dataset in machine learning and is often used for testing classification algorithms. It consists of 150 samples of iris flowers, each with four features (sepal length, sepal width, petal length, and petal width) and a target variable (the species of the flower).

CODE & OUTPUT:

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** jupyter Lab2_DataPreprocessingandDataVisualisationIris Last Checkpoint: 10 minutes ago (autosaved)
- Toolbar:** File, Edit, View, Insert, Cell, Kernel, Help, Not Trusted, Python 3 C
- In [1]:**

```
1 #####  
2 # LAB 2 - Iris Dataset EDA and Visualisation #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####
```
- In [2]:**

```
1 # Importing the Libraries  
2 import pandas as pd  
3 import matplotlib.pyplot as plt  
4 import numpy as np  
5 import seaborn as sns  
6  
7 #Additional libraries and commands  
8 import warnings # current version of seaborn generates a lot of warnings that can be ignored  
9 warnings.filterwarnings("ignore")
```
- In [3]:**

```
1 #Performing data preprocessing steps on Iris Dataset  
2 #1. Importing the dataset  
3 #2. Taking care of missing data, if any.  
4 #3. Encoding Categorical data, if required.  
5 #4. Train test splitting, if required.  
6 #5. Feature Scaling, if required.  
7  
8 #Since in this Lab, we are doing only EDA and visualisation, we can ignore steps 3, 4 and 5
```
- In [4]:**

```
1 #STEP 1 - Importing the Dataset  
2 iris = pd.read_csv("iris.csv")  
3 iris.head()
```
- Out[4]:**

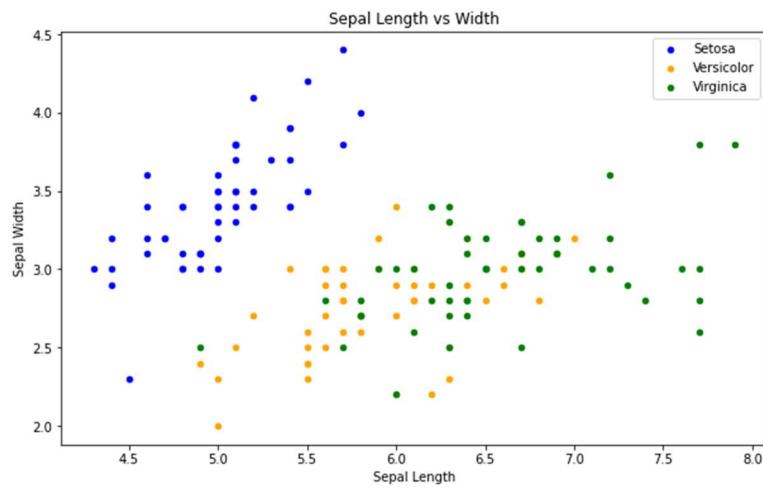
	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
In [5]: 1 #STEP 2 - Taking care of missing data, if any.  
2 iris.isna().sum()
```

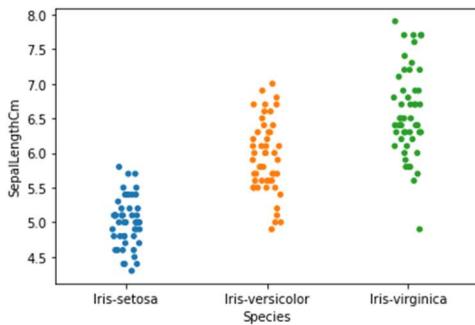
```
Out[5]: Id      0  
SepalLengthCm  0  
SepalWidthCm   0  
PetalLengthCm  0  
PetalWidthCm   0  
Species        0  
dtype: int64
```

```
In [6]: 1 #As we can see, there are no missing values.  
2 #Thus, we can see this dataset does not require much preprocessing. Which is to be expected considering  
3 #Iris dataset is one of the most simplest and commonly used small dataset.
```

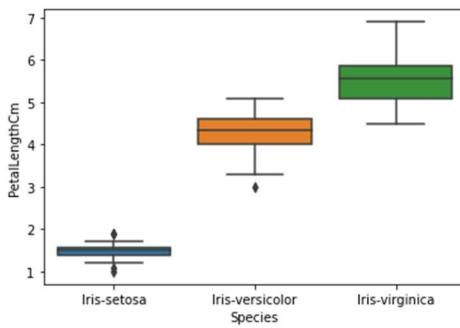
```
In [7]: 1 #Data Visualisation - Graphical visualisation.  
2  
3 #Using ScatterPlots of the Iris features  
4 fig = iris[iris.Species=='Iris-setosa'].plot(kind='scatter',x='SepalLengthCm',y='SepalWidthCm',color='blue',  
5                                label='Setosa', figsize=(10,6))  
6 iris[iris.Species=='Iris-versicolor'].plot(kind='scatter',x='SepalLengthCm',y='SepalWidthCm',color='orange',  
7                                label='Versicolor', ax=fig)  
8 iris[iris.Species=='Iris-virginica'].plot(kind='scatter',x='SepalLengthCm',y='SepalWidthCm',color='green',  
9                                label='Virginica', ax=fig)  
10 fig.set_xlabel("Sepal Length")  
11 fig.set_ylabel("Sepal Width")  
12 fig.set_title("Sepal Length vs Width")  
13 plt.show()
```



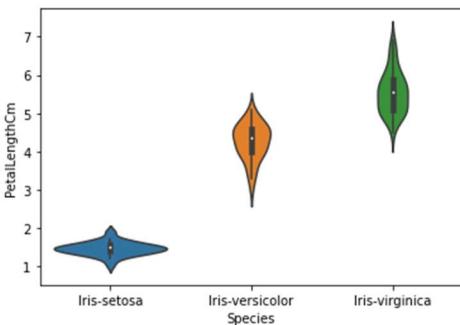
```
In [8]: 1 # Sepal Length using a Strip plot (Seaborn Library)
2 sns.stripplot(y = 'SepalLengthCm', x = 'Species', data =iris)
3 plt.show()
```



```
In [9]: 1 #Individual feature in through a boxplot (Seaborn)
2 sns.boxplot(x="Species", y="PetalLengthCm", data=iris)
3 plt.show()
```



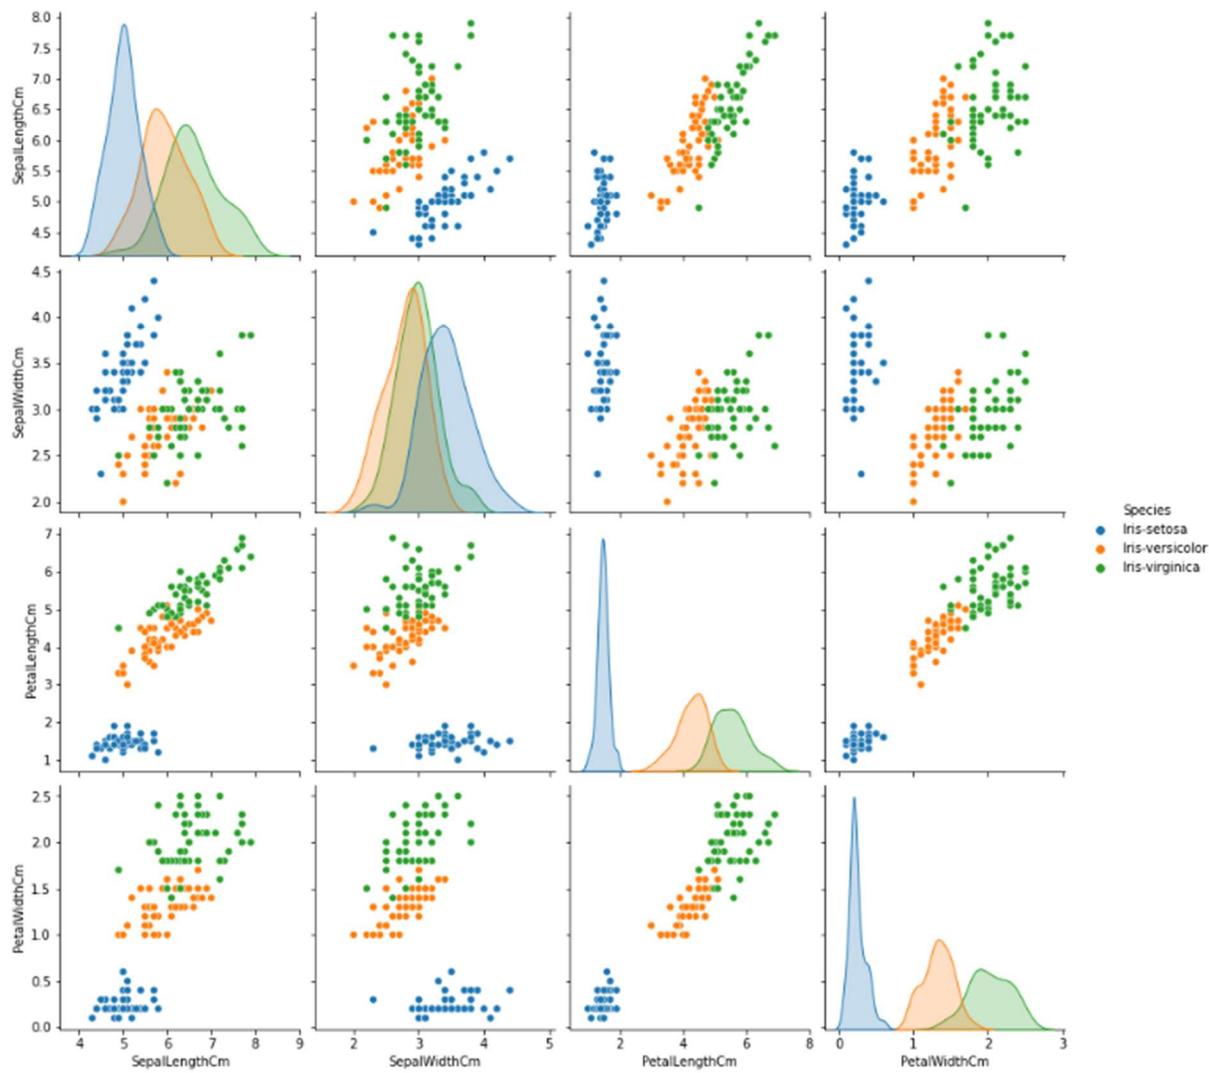
```
In [10]: 1 #Petal Length through Violin Plot
2 # A violin plot combines the benefits of the previous two plots (Strip Plot + Box Plot) and simplifies them
3 # Denser regions of the data are fatter, and sparser thinner in a violin plot
4
5 sns.violinplot(x="Species", y="PetalLengthCm", data=iris, size=6)
6 plt.show()
```



```
In [11]: 1 iris.columns
```

```
Out[11]: Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',
   'Species'],
   dtype='object')
```

```
In [12]: 1 #pairplot, which shows the bivariate relation between each pair of features
2 sns.pairplot(data = iris[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm','Species']],
3               hue="Species", size=3)
4 #ignoring ID column as it does not contribute anything to the dataset for classification or related purposes.
5 plt.show()
6 #From the pairplot, it is evident that Iris-setosa species is separated from the other two across all feature combinations
```



LEARNING OUTCOMES:

- Learn how to perform data preprocessing techniques such as data cleaning, feature scaling, and feature extraction.
- Gain an understanding of data visualization techniques and how to use them to explore and communicate patterns in data.
- Understand how to use Python libraries such as Pandas, NumPy, Matplotlib, and Seaborn for data analysis and visualization.
- From the pair plot, it is evident that Iris-Setosa species is separated from the other two across all feature combinations.

EXPERIMENT 3

AIM: To implement uni-variate linear regression using gradient descent.

THEORY:

Linear regression is a supervised machine learning algorithm used for predicting a continuous output variable based on one or more input variables. Uni-variate linear regression is a specific type of linear regression that involves predicting a single output variable based on a single input variable.

The main steps of the uni-variate linear regression using gradient descent algorithm are as follows:

1. Given a training dataset of input-output pairs, initialize the model parameters (slope and intercept) to some random values.
2. Compute the predicted output for each input using the current model parameters.
3. Calculate the error between the predicted output and the actual output for each input.
4. Update the model parameters using the gradient of the cost function with respect to each parameter. The cost function is typically the mean squared error between the predicted output and the actual output.
5. Repeat steps 2-4 until the model converges or a maximum number of iterations is reached.

Gradient descent is an optimization algorithm that is used to minimize the cost function by iteratively adjusting the model parameters in the direction of steepest descent.

CODE & OUTPUT:

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** jupyter Lab3_UniVariate_LinearRegression_GradientDescent Last Checkpoint: a few seconds ago (autosaved)
- Toolbar:** File, Edit, View, Insert, Cell, Kernel, Help, Not Trusted, Python 3 C, Logout
- In [1]:** Code cell containing:

```
1 #####  
2 # LAB 3 - Linear Regression using Gradient Descent #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####
```
- In [2]:** Code cell containing:

```
1 #importing the required Libraries  
2 import pandas as pd  
3 import numpy as np  
4 import matplotlib.pyplot as plt  
5 import seaborn as sns
```
- In [3]:** Code cell containing:

```
1 #importing the dataset  
2 dataset = pd.read_csv("advertising.csv")  
3 dataset[["TV", "Sales"]].head()
```
- Out[3]:** Data preview cell showing the first 5 rows of the dataset:

	TV	Sales
0	230.1	22.1
1	44.5	10.4
2	17.2	12.0
3	151.5	16.5
4	180.8	17.9
- In [4]:** Code cell containing:

```
1 dataset[["TV", "Sales"]].shape
```
- Out[4]:** Data output cell showing the shape of the dataset:

```
(200, 2)
```

```
In [5]: 1 #keeping only the first column and target column
2
3 reg_data = dataset[["TV","Sales"]]
4
5 #Separating the columns based on independent and dependent variable.
6 x = reg_data['TV']
7 y = reg_data['Sales']
8
9 #Applying standardisation to the independent variable
10 x = (x - x.mean()) / x.std()
11
12 #adding a column of ones to signify the bias to be added when we do dot product with theta vector
13 x = np.c_[np.ones(x.shape[0]), x]
```

```
In [6]: 1 x[:5]
```

```
Out[6]: array([[ 1.          ,  0.9674246 ],
   [ 1.          , -1.19437904],
   [ 1.          , -1.51235985],
   [ 1.          ,  0.05191939],
   [ 1.          ,  0.39319551]])
```

```
In [7]: 1 #Starting the gradient descent
2
3 #Setting the Learning rate
4 alpha = 0.01
5
6 #Setting the number of iterations
7 iterations = 2000
8
9 #Finding number of points available in the dataset - 200
10 m = y.size
11
12 #Setting seed so that our values can be replicated
13 np.random.seed(123)
14 #Assigning random values as the start to our gradient descent
15 theta = np.random.rand(2)
```

```
In [8]: 1 theta
```

```
Out[8]: array([0.69646919, 0.28613933])
```

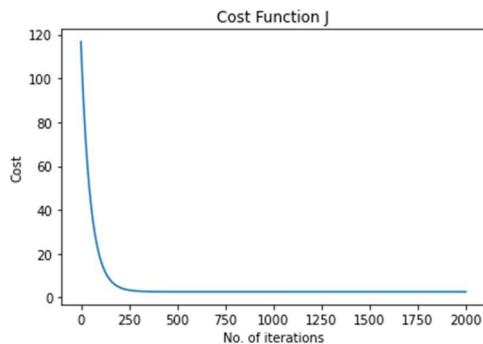
```
In [9]: 1 #Defining a function which apply gradient descent for the specified iterations and
2 #Calculate the final theta values
3 def gradient_descent(x, y, theta, iterations, alpha):
4
5     #Creating past records to help us plot the graph later on and evaluate the performance of our algo
6     past_costs = []
7     past_thetas = [theta]
8
9     #Actual loop to calculate predictions, then errors and cost, then gradients and finally update the thetas/weights
10    for i in range(iterations):
11        prediction = np.dot(x, theta)
12        error = prediction - y
13        cost = 1/(2*m) * np.dot(error.T, error)
14        past_costs.append(cost)
15        theta = theta - (alpha * (1/m) * np.dot(x.T, error))
16        past_thetas.append(theta)
17
18    return past_thetas, past_costs
```

```
In [10]: 1 #Actually running the above function on our dataset
2 past_thetas, past_costs = gradient_descent(x, y, theta, iterations, alpha)
3 theta = past_thetas[-1]
```

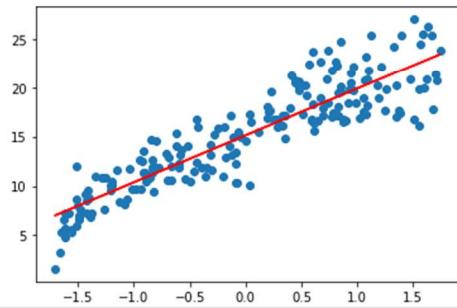
```
In [11]: 1 #Final values of our thetas
2 print(f" Gradient Descent: {round(theta[0],2)}, {round(theta[1],2)}")
```

```
Gradient Descent: 15.13, 4.76
```

```
In [12]: 1 #Plot the cost function...
2 plt.plot(past_costs)
3 plt.title('Cost Function J')
4 plt.xlabel('No. of iterations')
5 plt.ylabel('Cost')
6 plt.show()
```



```
In [13]: 1 #plot to show final Linear regression line
2 prediction = np.dot(x, theta)
3 plt.scatter(x[:,1],y)
4 plt.plot(x[:,1],prediction,color="red")
5 plt.show()
```



LEARNING OUTCOMES:

- Learn how to implement uni-variate linear regression using gradient descent.
- Understand the concept of a cost function and how to minimize it using gradient descent.
- Learn how to use NumPy and Matplotlib to generate random data, manipulate arrays, and visualize the data and model.

EXPERIMENT 4

AIM: To implement multi-variate linear regression using gradient descent.

THEORY:

Multi-variate linear regression is a supervised machine learning algorithm used for predicting a continuous output variable based on multiple input variables. In contrast to uni-variate linear regression, which involves predicting a single output variable based on a single input variable, multi-variate linear regression involves predicting a single output variable based on multiple input variables.

The main steps of the multi-variate linear regression using gradient descent algorithm are as follows:

1. Given a training dataset of input-output pairs, initialize the model parameters (slopes and intercept) to some random values.
2. Compute the predicted output for each input using the current model parameters.
3. Calculate the error between the predicted output and the actual output for each input.
4. Update the model parameters using the gradient of the cost function with respect to each parameter. The cost function is typically the mean squared error between the predicted output and the actual output.
5. Repeat steps 2-4 until the model converges or a maximum number of iterations is reached.

Gradient descent is an optimization algorithm that is used to minimize the cost function by iteratively adjusting the model parameters in the direction of steepest descent.

CODE & OUTPUT:

The screenshot shows a Jupyter Notebook interface with the following content:

- In [1]:** Python code for initializing parameters:

```
1 #####  
2 # LAB 4 - MultiVariate Linear Regression using GD #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####
```
- In [3]:** Python code for importing libraries:

```
1 #importing the required Libraries  
2 import pandas as pd  
3 import numpy as np  
4 import matplotlib.pyplot as plt
```
- In [4]:** Python code for importing the dataset:

```
1 #importing the dataset  
2 dataset = pd.read_csv("advertising.csv")  
3 dataset.head()
```
- Out[4]:** Data preview of the dataset:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
- In [5]:** Python code for checking dataset shape:

```
1 dataset.shape
```
- Out[5]:** Output of dataset shape:

```
(200, 4)
```
- In [6]:** Python code for checking dataset columns:

```
1 dataset.columns
```
- Out[6]:** Output of dataset columns:

```
Index(['TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')
```

```
In [11]: 1 #keeping only the first column and target column
2
3 reg_data = dataset.copy()
4
5 #Separating the columns based on independent and dependent variable.
6 x = reg_data[['TV', 'Radio', 'Newspaper']]
7 y = reg_data['Sales']
```

```
In [13]: 1 #Applying standardisation to the independent variable
2 x = (x - x.mean()) / x.std()
3
4 #adding a column of ones to signify the bias to be added when we do dot product with theta vector
5 x = np.c_[np.ones(x.shape[0]), x]
```

```
In [14]: 1 x[:5]
```

```
Out[14]: array([[ 1.          ,  0.9674246 ,  0.97906559,  1.77449253],
   [ 1.          , -1.19437904,  1.0800974 ,  0.66790272],
   [ 1.          , -1.51235985,  1.52463736,  1.77908419],
   [ 1.          ,  0.05191939,  1.21480648,  1.28318502],
   [ 1.          ,  0.39319551, -0.83950698,  1.27859336]])
```

```
In [17]: 1 #Starting the gradient descent
2
3 #Setting the Learning rate
4 alpha = 0.01
5
6 #Setting the number of iterations
7 iterations = 2000
8
9 #Finding number of points available in the dataset - 200
10 m = y.size
11
12 #Setting seed so that our values can be replicated
13 np.random.seed(123)
14 #Assigning random values as the start to our gradient descent
15 theta = np.random.rand(x.shape[1])
```

```
In [18]: 1 theta
```

```
Out[18]: array([0.69646919, 0.28613933, 0.22685145, 0.55131477])
```

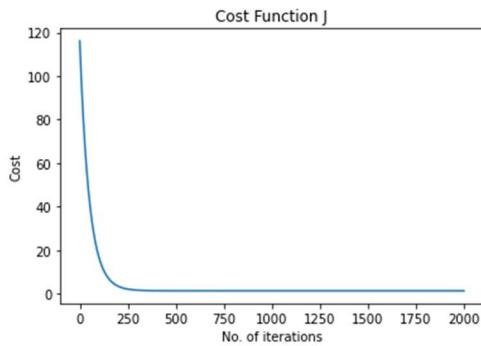
```
In [19]: 1 #Defining a function which apply gradient descent for the specified iterations and
2 #Calculate the final theta values
3 def gradient_descent(x, y, theta, iterations, alpha):
4
5     #Creating past records to help us plot the graph later on and evaluate the performance of our algo
6     past_costs = []
7     past_thetas = [theta]
8
9     #Actual Loop to calculate predictions, then errors and cost, then gradients and finally update the thetas/weights
10    for i in range(iterations):
11        prediction = np.dot(x, theta)
12        error = prediction - y
13        cost = 1/(2*m) * np.dot(error.T, error)
14        past_costs.append(cost)
15        theta = theta - (alpha * (1/m) * np.dot(x.T, error))
16        past_thetas.append(theta)
17
18    return past_thetas, past_costs
```

```
In [20]: 1 #Actually running the above function on our dataset
2 past_thetas, past_costs = gradient_descent(x, y, theta, iterations, alpha)
3 theta = past_thetas[-1]
```

```
In [26]: 1 #Final values of our thetas
2 print(f" Gradient Descent: {round(theta[0],2)}, {round(theta[1],2)}, {round(theta[2],2)}, {round(theta[3],2)}")
```

Gradient Descent: 15.13, 4.67, 1.59, 0.01

```
In [27]: 1 #Plot the cost function...
2 plt.plot(past_costs)
3 plt.title('Cost Function J')
4 plt.xlabel('No. of iterations')
5 plt.ylabel('Cost')
6 plt.show()
```



LEARNING OUTCOMES:

- Learn how to implement multi-variate linear regression using gradient descent.
- Understand the concept of a cost function and how to minimize it using gradient descent.
- Learn how to use NumPy and Matplotlib to generate random data, manipulate arrays, and visualize the data and model.

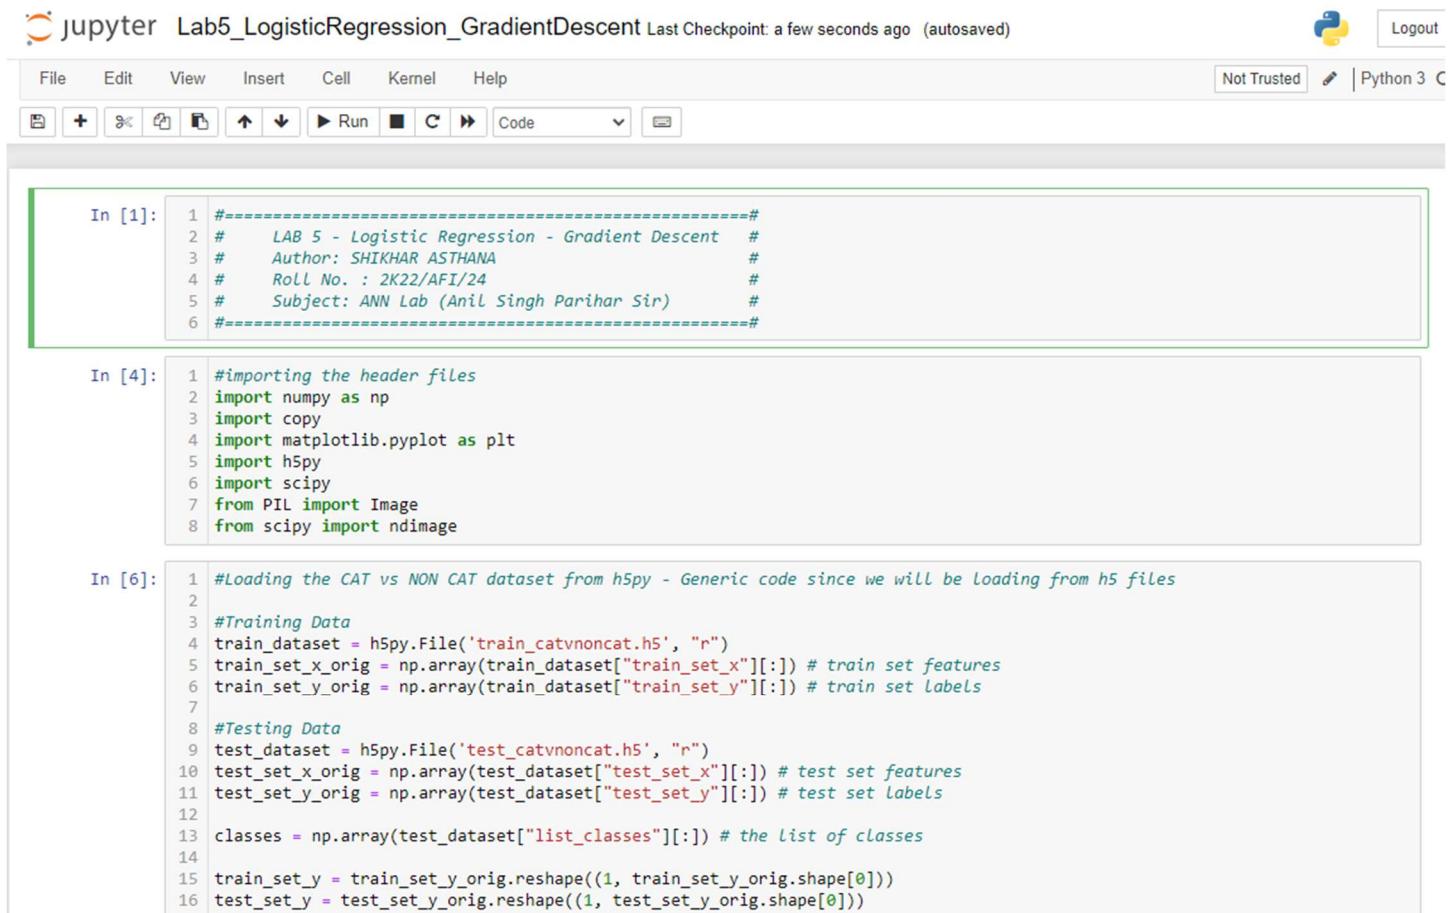
EXPERIMENT 5

AIM: To implement logistic regression using gradient descent.

THEORY:

- Logistic regression is a machine learning algorithm used for binary classification. It models the probability of a binary output variable (0 or 1) based on one or more input variables. It uses the sigmoid function to map any real-valued input to a value between 0 and 1.
- Gradient descent is an optimization algorithm used for finding the values of the parameters of a model that minimize a cost function. In logistic regression, the cost function is typically the log loss (also known as the cross-entropy loss).
- The gradient of a function is a vector that points in the direction of the steepest increase in the function. By taking steps in the opposite direction of the gradient, the parameters of a model can be updated to minimize the cost function.

CODE & OUTPUT:

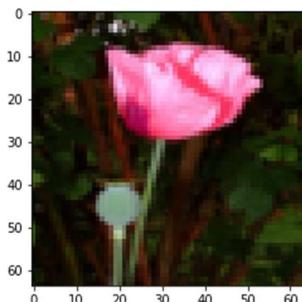


The screenshot shows a Jupyter Notebook interface with three code cells. The top cell (In [1]) contains a header comment. The second cell (In [4]) imports various Python libraries. The third cell (In [6]) loads a dataset from h5 files and performs some initial processing.

```
In [1]:  
1 #####  
2 # LAB 5 - Logistic Regression - Gradient Descent #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####  
  
In [4]:  
1 #importing the header files  
2 import numpy as np  
3 import copy  
4 import matplotlib.pyplot as plt  
5 import h5py  
6 import scipy  
7 from PIL import Image  
8 from scipy import ndimage  
  
In [6]:  
1 #Loading the CAT vs NON CAT dataset from h5py - Generic code since we will be Loading from h5 files  
2  
3 #Training Data  
4 train_dataset = h5py.File('train_catvnoncat.h5', "r")  
5 train_set_x_orig = np.array(train_dataset["train_set_x"][:]) # train set features  
6 train_set_y_orig = np.array(train_dataset["train_set_y"][:]) # train set labels  
7  
8 #Testing Data  
9 test_dataset = h5py.File('test_catvnoncat.h5', "r")  
10 test_set_x_orig = np.array(test_dataset["test_set_x"][:]) # test set features  
11 test_set_y_orig = np.array(test_dataset["test_set_y"][:]) # test set labels  
12  
13 classes = np.array(test_dataset["list_classes"][:]) # the list of classes  
14  
15 train_set_y = train_set_y_orig.reshape((1, train_set_y_orig.shape[0]))  
16 test_set_y = test_set_y_orig.reshape((1, test_set_y_orig.shape[0]))
```

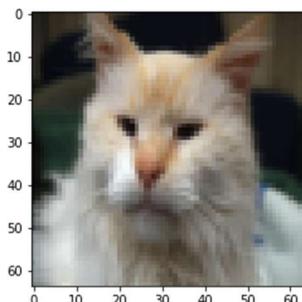
```
In [10]: 1 #checking a random image  
2 plt.imshow(train_set_x_orig[30])
```

```
Out[10]: <matplotlib.image.AxesImage at 0x19f206e5130>
```



```
In [9]: 1 plt.imshow(train_set_x_orig[27])
```

```
Out[9]: <matplotlib.image.AxesImage at 0x19f2067b160>
```



```
In [12]: 1 #Finding information about the number of training samples, number of testing samples and image size  
2  
3 m_train = train_set_x_orig.shape[0]  
4 m_test = test_set_x_orig.shape[0]  
5 num_px = train_set_x_orig.shape[1]  
6  
7 print ("Number of training examples: m_train = " + str(m_train))  
8 print ("Number of testing examples: m_test = " + str(m_test))  
9 print ("Height/Width of each image: num_px = " + str(num_px))  
10  
11 #Square images so image size is num_px * num_px * 3 (rgb)  
12 print ("Each image is of size: (" + str(num_px) + ", " + str(num_px) + ", 3)")  
13  
14 print ("train_set_x shape: " + str(train_set_x_orig.shape))  
15 print ("train_set_y shape: " + str(train_set_y.shape))  
16 print ("test_set_x shape: " + str(test_set_x_orig.shape))  
17 print ("test_set_y shape: " + str(test_set_y.shape))
```

```
Number of training examples: m_train = 209  
Number of testing examples: m_test = 50  
Height/Width of each image: num_px = 64  
Each image is of size: (64, 64, 3)  
train_set_x shape: (209, 64, 64, 3)  
train_set_y shape: (1, 209)  
test_set_x shape: (50, 64, 64, 3)  
test_set_y shape: (1, 50)
```

```
In [15]: 1 #converting the images into a single flattened input vector  
2 train_set_x_flatten = train_set_x_orig.reshape(train_set_x_orig.shape[0], -1).T  
3 test_set_x_flatten = test_set_x_orig.reshape(test_set_x_orig.shape[0], -1).T  
4  
5 print(train_set_x_flatten.shape)  
6 print(test_set_x_flatten.shape)
```

```
(12288, 209)  
(12288, 50)
```

```
In [16]: 1 #We need to standardise the images so that the value is 0-255  
2 train_set_x = train_set_x_flatten / 255.  
3 test_set_x = test_set_x_flatten / 255.
```

```
In [17]: 1 #Defining the sigmoid function which we will be using in our logistic regression
2 def sigmoid(z):
3     s = 1/(1+np.exp(-z))
4     return s
5
6 #Defining a function to initialise the weight vector with 0 values
7 def initialize_with_zeros(dim):
8     w = np.zeros((dim,1))
9     b = 0.0
10
11    return w, b
12
13
14 #Defining a function which will calculate the forward pass and backward pass of the algorithm once
15 def propagate(w, b, X, Y):
16     m = X.shape[1]
17
18     #Forward propagation
19     A = sigmoid(np.dot(w.T,X)+b)
20     cost = (-1/m)*np.sum(np.dot(np.log(A),Y.T)+np.dot(np.log(1-A),1-Y.T))
21
22     #backward pass
23     dw = (1/m)*np.dot(X,(A-Y).T)
24     db = (1/m)*np.sum(A-Y)
25
26     #cost will be an array which will retain costs of previous cycles as well
27     cost = np.squeeze(np.array(cost))
28
29     #grads will only be for the current cycle
30     grads = {"dw": dw,
31               "db": db}
32
33    return grads, cost
```

```
In [18]: 1 #Function to actually implement the gradient descent algorithm
2 def optimize(w, b, X, Y, num_iterations=100, learning_rate=0.009, print_cost=False):
3
4     #creating deep copies so that we can update them simultaneously at the end
5     w = copy.deepcopy(w)
6     b = copy.deepcopy(b)
7
8     costs = []
9
10    for i in range(num_iterations):
11        #Applying the forward and backward pass
12        grads,cost = propagate(w,b,X,Y)
13
14        # Retrieve derivatives from grads
15        dw = grads["dw"]
16        db = grads["db"]
17
18        #Updating weights using gradient descent
19        w = w - (learning_rate * dw)
20        b = b - (learning_rate * db)
21
22        # Record the costs
23        if i % 100 == 0:
24            costs.append(cost)
25
26        # Print the cost every 100 training iterations
27        if print_cost:
28            print ("Cost after iteration %i: %f" %(i, cost))
29
30
31    params = {"w": w,
32              "b": b}
33
34    grads = {"dw": dw,
35              "db": db}
36
37    return params, grads, costs
```

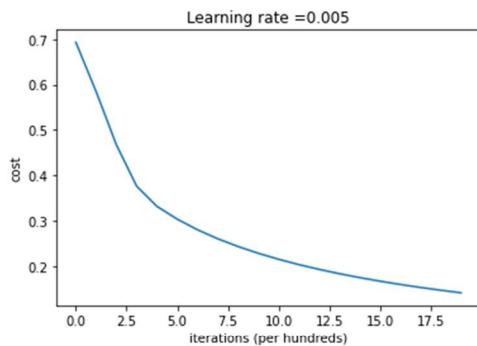
```
In [19]: 1 #Function to predict
2 def predict(w, b, X):
3
4     m = X.shape[1]
5     Y_prediction = np.zeros((1, m))
6     w = w.reshape(X.shape[0], 1)
7
8     #calculating the values using dot product
9     A = sigmoid(np.dot(w.T,X)+b)
10
11    #these values need to be assigned either 0 or 1 since logistic regression does not work on continuous values
12    for i in range(A.shape[1]):
13        if A[0,i] > 0.5:
14            Y_prediction[0,i] = 1
15        else:
16            Y_prediction[0,i] = 0
17
18    return Y_prediction
```

```
In [20]: 1 #Building a complete Logistic regression model using the above defined functions
2 def model(X_train, Y_train, X_test, Y_test, num_iterations=2000, learning_rate=0.5, print_cost=False):
3
4     #Initialise the weights
5     w,b = initialize_with_zeros(X_train.shape[0])
6
7     #Apply gradient descent
8     params, grads, costs = optimize(w, b, X_train, Y_train, num_iterations, learning_rate, print_cost)
9     w = params["w"]
10    b = params["b"]
11
12    #Get predictions for both train and test
13    Y_prediction_test = predict(w, b, X_test)
14    Y_prediction_train = predict(w, b, X_train)
15
16    # Print train/test Errors
17    if print_cost:
18        print("train accuracy: {} %".format(100 - np.mean(np.abs(Y_prediction_train - Y_train)) * 100))
19        print("test accuracy: {} %".format(100 - np.mean(np.abs(Y_prediction_test - Y_test)) * 100))
20
21
22    d = {"costs": costs,
23          "Y_prediction_test": Y_prediction_test,
24          "Y_prediction_train" : Y_prediction_train,
25          "w" : w,
26          "b" : b,
27          "learning_rate" : learning_rate,
28          "num_iterations": num_iterations}
29
30    return d
```

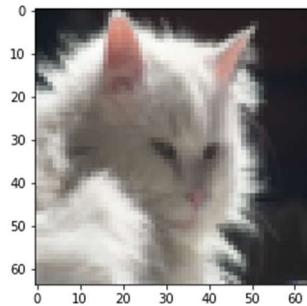
```
In [21]: 1 #actually running the aove defined model on our cats dataset
2 logistic_regression_model = model(train_set_X, train_set_y, test_set_X, test_set_y, num_iterations=2000,
3                                     learning_rate=0.005, print_cost=True)
4
```

```
Cost after iteration 0: 0.693147
Cost after iteration 100: 0.584508
Cost after iteration 200: 0.466949
Cost after iteration 300: 0.376007
Cost after iteration 400: 0.331463
Cost after iteration 500: 0.303273
Cost after iteration 600: 0.279880
Cost after iteration 700: 0.260042
Cost after iteration 800: 0.242941
Cost after iteration 900: 0.228004
Cost after iteration 1000: 0.214820
Cost after iteration 1100: 0.203078
Cost after iteration 1200: 0.192544
Cost after iteration 1300: 0.183033
Cost after iteration 1400: 0.174399
Cost after iteration 1500: 0.166521
Cost after iteration 1600: 0.159305
Cost after iteration 1700: 0.152667
Cost after iteration 1800: 0.146542
Cost after iteration 1900: 0.140872
train accuracy: 99.04306220095694 %
test accuracy: 70.0 %
```

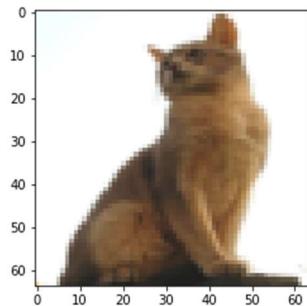
```
In [22]: 1 # Plot Learning curve (with costs)
2 costs = np.squeeze(logistic_regression_model['costs'])
3 plt.plot(costs)
4 plt.ylabel('cost')
5 plt.xlabel('iterations (per hundreds)')
6 plt.title("Learning rate = " + str(logistic_regression_model["learning_rate"]))
7 plt.show()
```



```
In [25]: 1 #checking random output
2 index = 8
3 plt.imshow(test_set_x[:, index].reshape((num_px, num_px, 3)))
4 print ("y = " + str(test_set_y[0,index]) + ", Prediction: " + classes[int(logistic_regression_model['Y_prediction_test'])[0,i]] )
y = 1, Prediction: cat
```



```
In [27]: 1 #checking random output
2 index = 6
3 plt.imshow(test_set_x[:, index].reshape((num_px, num_px, 3)))
4 print ("y = " + str(test_set_y[0,index]) + ", Prediction: " + classes[int(logistic_regression_model['Y_prediction_test'])[0,i]] )
y = 1, Prediction: non-cat
```



LEARNING OUTCOMES:

- Learn how to implement logistic regression using gradient descent.
- Understand the concept of a cost function and how to minimize it using gradient descent.
- Learn how to use NumPy and Matplotlib to generate random data, manipulate arrays, and visualize the data and model.

EXPERIMENT 6

AIM: To implement basic perceptron model.

THEORY:

Perceptron is a simple supervised learning algorithm used for binary classification. It takes in input values and produces a binary output (0 or 1) based on a weighted sum of those input values. The weights are adjusted during training in order to minimize the classification error.

The basic steps of the perceptron algorithm are as follows:

1. Initialize the weights to small random values.
2. For each input-output pair in the training data:
 - a. Calculate the weighted sum of the inputs.
 - b. Apply an activation function to the weighted sum to obtain the predicted output.
 - c. Update the weights based on the difference between the predicted output and the true output.
3. Repeat step 2 for a fixed number of epochs or until convergence.

The perceptron algorithm is a linear classification algorithm, which means it can only separate data that is linearly separable.

CODE & OUTPUT:

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** jupyter Lab6_BasicPerceptronModel Last Checkpoint: 3 minutes ago (unsaved changes)
- Toolbar:** File, Edit, View, Insert, Cell, Kernel, Help, Trusted, Python 3 C
- In [1]:**

```
1 #####  
2 # LAB 6 - Implement Basic Perceptron Model #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####
```
- In [2]:**

```
1 #Importing header files  
2 import pandas as pd  
3 import numpy as np
```
- In [3]:**

```
1 #For input, I will make a dataframe for OR Gate  
2 #This can be changed to AND or any other Linearly separable combination of inputs  
3 input_data = pd.DataFrame({  
4     'x1': [0, 0, 1, 1],  
5     'x2': [0, 1, 0, 1],  
6     'y': [0, 1, 1, 1]  
7 })
```
- In [4]:**

```
1 #Defining the function for the perceptron  
2 def perceptron(inputs, weights):  
3     # calculate the weighted sum of the inputs and weights  
4     weighted_sum = np.dot(inputs, weights)  
5  
6     # apply the step function to the weighted sum  
7     if weighted_sum > 0:  
8         return 1  
9     else:  
10        return 0
```

```
In [5]: 1 # set the learning rate and maximum number of iterations
2 learning_rate = 0.1
3 max_iterations = 100
4
5 # randomly initialize the weights
6 weights = np.random.uniform(size=3)
7
8 # iterate through the data and update the weights
9 for i in range(max_iterations):
10    # iterate through each row in the input data
11    for index, row in input_data.iterrows():
12        # extract the inputs and output for this row
13        inputs = row[['x1', 'x2']]
14        expected_output = row['y']
15
16        # add a bias term to the inputs
17        inputs_with_bias = np.append(inputs, 1)
18
19        # calculate the predicted output using the perceptron function
20        predicted_output = perceptron(inputs_with_bias, weights)
21
22        # calculate the error and update the weights
23        error = expected_output - predicted_output
24        weights += learning_rate * error * inputs_with_bias
25
26        # check if the model has converged
27        if error == 0:
28            break
```

```
In [6]: 1 #Checking final weights
2 print(weights)
[ 0.23507456  0.53046904 -0.00477351]
```

```
In [7]: 1 # create a new input with values [1, 1]
2 new_input = np.array([1, 1, 1])
3
4 # predict the output using the perceptron
5 predicted_output = perceptron(new_input, weights)
6
7 print(predicted_output) # should output 1
```

1

LEARNING OUTCOMES:

- Understanding the basic concepts of supervised machine learning and binary classification tasks.
- Familiarity with the perceptron algorithm and its use for binary classification tasks.
- Ability to implement the perceptron algorithm using Python and the numpy library, including how to calculate the weighted sum, apply the activation function, and update the weights.
- Understanding of the limitations of the perceptron algorithm, including its inability to handle non-linearly separable data and its sensitivity to the initial weights.
- Familiarity with techniques for improving the performance of the model, such as using a different activation function or incorporating regularization.

EXPERIMENT 7

AIM: To implement multi-layer neural network.

THEORY:

A multi-layer neural network is a type of deep learning model that consists of multiple layers of interconnected nodes (neurons). Each layer performs a set of computations on its inputs and produces an output that serves as the input to the next layer. The first layer is called the input layer, the last layer is called the output layer, and all other layers are called hidden layers.

The basic steps of training a multi-layer neural network are as follows:

1. Initialize the weights and biases to small random values.
2. For each input-output pair in the training data:
 - a. Feed the input forward through the network to obtain the predicted output.
 - b. Calculate the error between the predicted output and the true output.
 - c. Backpropagate the error through the network to update the weights and biases.
3. Repeat step 2 for a fixed number of epochs or until convergence.

During training, the weights and biases are updated using an optimization algorithm such as stochastic gradient descent. There are also many variations of the basic neural network architecture, including different activation functions, regularization techniques, and architectures such as convolutional neural networks and recurrent neural networks.

The Boston Housing Dataset is a derived from information collected by the U.S. Census Service concerning housing in the area of Boston MA. The following describes the dataset columns:

1. CRIM - per capita crime rate by town
2. ZN - proportion of residential land zoned for lots over 25,000 sq.ft.
3. INDUS - proportion of non-retail business acres per town.
4. CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise)
5. NOX - nitric oxides concentration (parts per 10 million)
6. RM - average number of rooms per dwelling
7. AGE - proportion of owner-occupied units built prior to 1940
8. DIS - weighted distances to five Boston employment centers
9. RAD - index of accessibility to radial highways
10. TAX - full-value property-tax rate per \$10,000
11. PTRATIO - pupil-teacher ratio by town
12. B - $1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town
13. LSTAT - % lower status of the population
14. MEDV - Median value of owner-occupied homes in \$1000's

Variables 1-13 will be our independent variables.

Variable 14 would be our dependent variable.

CODE & OUTPUT:

jupyter Lab7_MultiLayerNeuralNetwork_BostonHousing_Dataset Last Checkpoint: a few seconds ago (autosaved)  Logout

File Edit View Insert Cell Kernel Help Trusted Python 3 C

In [1]:

```
1 #####  
2 # LAB 7 - MultiLayer Neural Network #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####
```

In [17]:

```
1 #Importing header files  
2 import tensorflow as tf  
3 from tensorflow import keras  
4 from sklearn.datasets import load_boston  
5 from sklearn.model_selection import train_test_split  
6 from sklearn.preprocessing import StandardScaler  
7 import matplotlib.pyplot as plt
```

In [4]:

```
1 #Using the shallow neural network for house price prediction  
2 #Using the prebuild dataset - Boston Housing Dataset  
3 boston = load_boston()  
4 X = boston.data  
5 y = boston.target
```

In [9]:

```
1 X[:1][0]
```

Out[9]:

```
array([6.320e-03, 1.800e+01, 2.310e+00, 0.000e+00, 5.380e-01, 6.575e+00,  
       6.520e+01, 4.090e+00, 1.000e+00, 2.960e+02, 1.530e+01, 3.969e+02,  
       4.980e+00])
```

In [10]:

```
1 y[:1][0]
```

Out[10]:

```
24.0
```

In [11]:

```
1 #Splitting the dataset into train test  
2 #Using 80 - 20 split  
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

In [12]:

```
1 #Applying standardization for the datasets  
2 scaler = StandardScaler()  
3 X_train = scaler.fit_transform(X_train)  
4 X_test = scaler.transform(X_test)
```

In [13]:

```
1 #Building a sequential neural network model  
2 #Using one hidden layer containing 32 neurons  
3 #A ReLU activation function  
4 #Linear activation function in the output layer, as this is a regression problem  
5 model = keras.Sequential([  
6     keras.layers.Dense(32, activation='relu', input_shape=(X_train.shape[1],)),  
7     keras.layers.Dropout(0.5),  
8     keras.layers.Dense(1, activation='linear')  
9 ])
```

In [14]:

```
1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 32)	448
dropout (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 1)	33
Total params:	481	
Trainable params:	481	
Non-trainable params:	0	

In [15]:

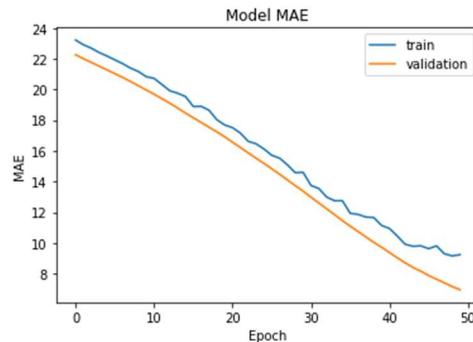
```
1 #Defining generic Loss functions, optimizers and metrics to be used by our model in training  
2 model.compile(loss='mean_squared_error', optimizer='adam', metrics=['mean_absolute_error'])
```

```
In [16]: 1 #Actually training the model
2 history = model.fit(X_train, y_train, epochs=50, batch_size=32, validation_split=0.2)
3
4 test_loss, test_mae = model.evaluate(X_test, y_test)
5 print('Test MAE:', test_mae)

323/323 [=====] - 0s 95us/sample - loss: 144.0639 - mean_absolute_error: 9.8165 - val_loss: 99.5749
- val_mean_absolute_error: 8.1632
Epoch 46/50
323/323 [=====] - 0s 142us/sample - loss: 138.1007 - mean_absolute_error: 9.6319 - val_loss: 94.1519
- val_mean_absolute_error: 7.8846
Epoch 47/50
323/323 [=====] - 0s 88us/sample - loss: 146.8096 - mean_absolute_error: 9.8130 - val_loss: 89.8243
- val_mean_absolute_error: 7.6433

Epoch 48/50
323/323 [=====] - 0s 94us/sample - loss: 136.8551 - mean_absolute_error: 9.2996 - val_loss: 85.7512
- val_mean_absolute_error: 7.4069
Epoch 49/50
323/323 [=====] - 0s 91us/sample - loss: 124.8400 - mean_absolute_error: 9.1616 - val_loss: 81.7915
- val_mean_absolute_error: 7.1644
Epoch 50/50
323/323 [=====] - 0s 88us/sample - loss: 129.9768 - mean_absolute_error: 9.2384 - val_loss: 78.3691
- val_mean_absolute_error: 6.9540
102/102 [=====] - 0s 65us/sample - loss: 71.3065 - mean_absolute_error: 6.8688
Test MAE: 6.868761
```

```
In [18]: 1 # Plot training history
2 plt.plot(history.history['mean_absolute_error'])
3 plt.plot(history.history['val_mean_absolute_error'])
4 plt.title('Model MAE')
5 plt.ylabel('MAE')
6 plt.xlabel('Epoch')
7 plt.legend(['train', 'validation'], loc='upper right')
8 plt.show()
```



LEARNING OUTCOMES:

- Understanding the basic concepts of deep learning and neural networks.
- Familiarity with the architecture of a multi-layer neural network, including input, output, and hidden layers.
- Understanding of the training process for a neural network, including forward propagation, backpropagation, and weight updates.
- Familiarity with techniques for improving the performance of a neural network, such as using different activation functions or regularization techniques.
- Knowledge of the various applications of neural networks, such as image classification, natural language processing, and time series analysis.

EXPERIMENT 8

AIM: To perform classification on MNIST Digit dataset using Neural Network.

THEORY:

The MNIST dataset is a large collection of handwritten digit images that is commonly used as a benchmark dataset for image classification tasks. The dataset consists of 60,000 training images and 10,000 testing images, with each image being a grayscale image of size 28x28 pixels.

Neural networks are a class of machine learning algorithms that are designed to model the behavior of the human brain. They consist of interconnected nodes (neurons) that process information and produce an output. Neural networks can be used for a variety of tasks, including image classification.

To perform classification on the MNIST dataset using a neural network, we typically follow these steps:

1. Preprocess the data: normalize the pixel values, and split the data into training and testing sets.
2. Define the neural network architecture: specify the number and size of the layers, the activation functions, and the loss function.
3. Train the network: feed the training data through the network and adjust the weights and biases using an optimization algorithm such as stochastic gradient descent.
4. Evaluate the network: feed the testing data through the network and calculate the accuracy of the predictions.

CODE & OUTPUT:

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** jupyter Lab8_MNIST_Digit_Classification Last Checkpoint: a few seconds ago (autosaved)
- Toolbar:** File, Edit, View, Insert, Cell, Kernel, Help, Run, Cell Type (Code), Trusted, Python 3 C, Logout
- In [1]:** Python code for header comments:

```
1 #####  
2 # LAB 8-Classification MNIST using Neural Network #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####
```
- In [2]:** Python code for importing libraries:

```
1 #Import the header files  
2 import tensorflow as tf  
3 from tensorflow import keras  
4 from sklearn.model_selection import train_test_split  
5 from sklearn.preprocessing import StandardScaler  
6 from tensorflow.keras.datasets import mnist
```
- In [3]:** Python code for checking GPU compatibility:

```
1 #Checking if GPU is compatible  
2 physical_devices = tf.config.list_physical_devices('GPU')  
3 tf.config.experimental.set_memory_growth(physical_devices[0], True)
```
- In [4]:** Python code for loading the MNIST dataset:

```
1 #Loading MNIST Digit Dataset  
2 (X_train_full, y_train_full), (X_test, y_test) = mnist.load_data()  
3 X_train, X_val, y_train, y_val = train_test_split(X_train_full, y_train_full, test_size=0.2, random_state=42)
```
- In [5]:** Python code for reshaping data:

```
1 #Reshaping data so that every pixel value can be fed into the model  
2 X_train = X_train.reshape(X_train.shape[0], 784)  
3 X_val = X_val.reshape(X_val.shape[0], 784)  
4 X_test = X_test.reshape(X_test.shape[0], 784)
```
- In [6]:** Python code for scaling pixel values:

```
1 #Scaling the pixel values  
2 scaler = StandardScaler()  
3 X_train = scaler.fit_transform(X_train)  
4 X_val = scaler.transform(X_val)  
5 X_test = scaler.transform(X_test)
```

```
In [7]: 1 #Building the model
2 #Currently only 1 hidden layer
3 model = keras.Sequential([
4     keras.layers.Dense(128, activation='relu', input_shape=(784,)),
5     keras.layers.Dropout(0.2),
6     keras.layers.Dense(10, activation='softmax')
7 ])
```

```
In [8]: 1 model.summary()
```

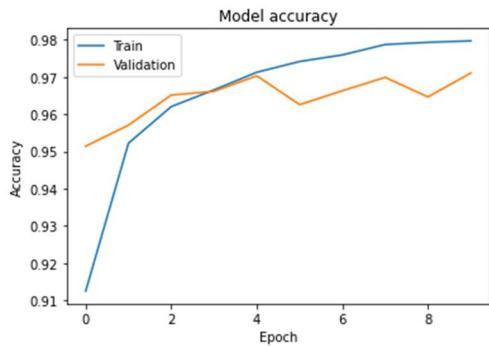
```
Model: "sequential"
-----  
Layer (type)      Output Shape       Param #
-----  
dense (Dense)    (None, 128)        100480  
dropout (Dropout) (None, 128)        0  
dense_1 (Dense)  (None, 10)         1290  
-----  
Total params: 101,770  
Trainable params: 101,770  
Non-trainable params: 0
```

```
In [9]: 1 model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [10]: 1 history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_val, y_val))
2
3 test_loss, test_acc = model.evaluate(X_test, y_test)
4 print('Test accuracy:', test_acc)
```

```
Train on 48000 samples, validate on 12000 samples
Epoch 1/10
48000/48000 [=====] - 85s 2ms/sample - loss: 0.3116 - accuracy: 0.9125 - val_loss: 0.2300 - val_accuracy: 0.9514
Epoch 2/10
48000/48000 [=====] - 2s 50us/sample - loss: 0.1759 - accuracy: 0.9522 - val_loss: 0.2219 - val_accuracy: 0.9571
Epoch 3/10
48000/48000 [=====] - 2s 49us/sample - loss: 0.1314 - accuracy: 0.9621 - val_loss: 0.2143 - val_accuracy: 0.9652
Epoch 4/10
48000/48000 [=====] - 2s 49us/sample - loss: 0.1155 - accuracy: 0.9666 - val_loss: 0.2018 - val_accuracy: 0.9662
Epoch 5/10
48000/48000 [=====] - 2s 48us/sample - loss: 0.0947 - accuracy: 0.9713 - val_loss: 0.2050 - val_accuracy: 0.9703
Epoch 6/10
48000/48000 [=====] - 3s 52us/sample - loss: 0.0855 - accuracy: 0.9742 - val_loss: 0.2848 - val_accuracy: 0.9626
Epoch 7/10
48000/48000 [=====] - 2s 49us/sample - loss: 0.0760 - accuracy: 0.9760 - val_loss: 0.3622 - val_accuracy: 0.9663
Epoch 8/10
48000/48000 [=====] - 2s 48us/sample - loss: 0.0677 - accuracy: 0.9787 - val_loss: 0.3407 - val_accuracy: 0.9699
Epoch 9/10
48000/48000 [=====] - 2s 49us/sample - loss: 0.0791 - accuracy: 0.9793 - val_loss: 0.3061 - val_accuracy: 0.9647
Epoch 10/10
48000/48000 [=====] - 2s 51us/sample - loss: 0.0686 - accuracy: 0.9797 - val_loss: 0.3045 - val_accuracy: 0.9711
10000/10000 [=====] - 0s 43us/sample - loss: 0.1843 - accuracy: 0.9708
Test accuracy: 0.9708
```

```
In [11]: 1 # Plot training history
2 import matplotlib.pyplot as plt
3
4 plt.plot(history.history['accuracy'])
5 plt.plot(history.history['val_accuracy'])
6 plt.title('Model accuracy')
7 plt.ylabel('Accuracy')
8 plt.xlabel('Epoch')
9 plt.legend(['Train', 'Validation'], loc='upper left')
10 plt.show()
```



LEARNING OUTCOMES:

- Understanding of the MNIST dataset and its use as a benchmark dataset for image classification tasks.
- Familiarity with the architecture of a neural network for image classification, including the number and size of layers, activation functions, and loss function.
- Ability to implement a neural network using Python and a deep learning library such as TensorFlow or PyTorch for image classification on the MNIST dataset.

EXPERIMENT 9

AIM: To perform classification on MNIST Fashion dataset using Neural Network.

THEORY:

The MNIST Fashion dataset is a collection of grayscale images of clothing items, such as t-shirts, dresses, shoes, and bags. The dataset consists of 60,000 training images and 10,000 testing images, each of size 28x28 pixels. The goal of the classification task is to classify each image into one of 10 categories representing different clothing types.

Neural networks are a class of machine learning algorithms that are designed to model the behavior of the human brain. They consist of interconnected nodes (neurons) that process information and produce an output. Neural networks can be used for a variety of tasks, including image classification.

To perform classification on the MNIST dataset using a neural network, we typically follow these steps:

1. Preprocess the data: normalize the pixel values, and split the data into training and testing sets.
2. Define the neural network architecture: specify the number and size of the layers, the activation functions, and the loss function.
3. Train the network: feed the training data through the network and adjust the weights and biases using an optimization algorithm such as stochastic gradient descent.
4. Evaluate the network: feed the testing data through the network and calculate the accuracy of the predictions.

CODE & OUTPUT:

The screenshot shows a Jupyter Notebook interface with several code cells and their corresponding outputs. The notebook title is "jupyter Lab9_MNIST_Fashion_Classification".

- In [1]:**

```
1 #####  
2 # LAB 9 - MNIST Fashion using Neural Network #  
3 # Author: SHIKHAR ASTHANA #  
4 # Roll No. : 2K22/AFI/24 #  
5 # Subject: ANN Lab (Anil Singh Parihar Sir) #  
6 #####
```
- In [2]:**

```
1 #Import the header files  
2 import tensorflow as tf  
3 from tensorflow import keras  
4 from sklearn.model_selection import train_test_split  
5 from sklearn.preprocessing import StandardScaler  
6 from tensorflow.keras.datasets import mnist
```
- In [3]:**

```
1 #Checking if GPU is compatible  
2 physical_devices = tf.config.list_physical_devices('GPU')  
3 tf.config.experimental.set_memory_growth(physical_devices[0], True)
```
- In [4]:**

```
1 #Loading MNIST Digit Dataset  
2 (X_train_full, y_train_full), (X_test, y_test) = tf.keras.datasets.fashion_mnist.load_data()  
3 X_train, X_val, y_train, y_val = train_test_split(X_train_full, y_train_full, test_size=0.2, random_state=42)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>
ERROR! Session/line number was not unique in database. History logging moved to new session 135
32768/29515 [=====] - 0s 1us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>
26427392/26421880 [=====] - 4s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>
8192/5148 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>
4423680/4422102 [=====] - 0s 0us/step
- In [5]:**

```
1 #Reshaping data so that every pixel value can be fed into the model  
2 X_train = X_train.reshape(X_train.shape[0], 784)  
3 X_val = X_val.reshape(X_val.shape[0], 784)  
4 X_test = X_test.reshape(X_test.shape[0], 784)
```

```
In [6]: 1 #Scaling the pixel values
2 scaler = StandardScaler()
3 X_train = scaler.fit_transform(X_train)
4 X_val = scaler.transform(X_val)
5 X_test = scaler.transform(X_test)
```

```
In [7]: 1 #Building the model
2 #Currently only 3 hidden layer
3 model = keras.Sequential([
4     keras.layers.Dense(128, activation='relu', input_shape=(784,)),
5     keras.layers.Dense(64, activation='relu'),
6     keras.layers.Dense(32, activation='relu'),
7     keras.layers.Dropout(0.2),
8     keras.layers.Dense(10, activation='softmax')
9 ])
```

```
In [8]: 1 model.summary()
```

```
Model: "sequential"
-----  
Layer (type)          Output Shape         Param #
-----  
dense (Dense)        (None, 128)          100480  
dense_1 (Dense)      (None, 64)           8256  
dense_2 (Dense)      (None, 32)           2080  
dropout (Dropout)    (None, 32)           0  
dense_3 (Dense)      (None, 10)            330  
-----  
Total params: 111,146
Trainable params: 111,146
Non-trainable params: 0
```

```
In [9]: 1 model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [10]: 1 history = model.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_val, y_val))
2
3 test_loss, test_acc = model.evaluate(X_test, y_test)
4 print('Test accuracy:', test_acc)
```

```
Train on 48000 samples, validate on 12000 samples
Epoch 1/20
48000/48000 [=====] - 100s 2ms/sample - loss: 0.5328 - accuracy: 0.8128 - val_loss: 0.3998 - val_accuracy: 0.8543
Epoch 2/20
48000/48000 [=====] - 5s 112us/sample - loss: 0.3876 - accuracy: 0.8622 - val_loss: 0.3563 - val_accuracy: 0.8705
Epoch 3/20
48000/48000 [=====] - 5s 113us/sample - loss: 0.3480 - accuracy: 0.8744 - val_loss: 0.3567 - val_accuracy: 0.8730
Epoch 4/20
48000/48000 [=====] - 6s 115us/sample - loss: 0.3237 - accuracy: 0.8824 - val_loss: 0.3576 - val_accuracy: 0.8732
Epoch 5/20
48000/48000 [=====] - 6s 116us/sample - loss: 0.3001 - accuracy: 0.8897 - val_loss: 0.3425 - val_accuracy: 0.8804
Epoch 6/20
48000/48000 [=====] - 5s 113us/sample - loss: 0.2811 - accuracy: 0.8969 - val_loss: 0.3378 - val_accuracy: 0.8800
Epoch 7/20
48000/48000 [=====] - 6s 115us/sample - loss: 0.2713 - accuracy: 0.9001 - val_loss: 0.3554 - val_accuracy: 0.8780
Epoch 8/20
48000/48000 [=====] - 6s 116us/sample - loss: 0.2540 - accuracy: 0.9065 - val_loss: 0.3341 - val_accuracy: 0.8840
Epoch 9/20
48000/48000 [=====] - 5s 114us/sample - loss: 0.2433 - accuracy: 0.9107 - val_loss: 0.3656 - val_accuracy: 0.8814
Epoch 10/20
48000/48000 [=====] - 5s 111us/sample - loss: 0.2340 - accuracy: 0.9134 - val_loss: 0.3557 - val_accuracy: 0.8833
Epoch 11/20
48000/48000 [=====] - 6s 118us/sample - loss: 0.2229 - accuracy: 0.9165 - val_loss: 0.3575 - val_accuracy: 0.8856
Epoch 12/20
48000/48000 [=====] - 6s 122us/sample - loss: 0.2145 - accuracy: 0.9193 - val_loss: 0.3443 - val_accuracy: 0.8899
Epoch 13/20
48000/48000 [=====] - 6s 116us/sample - loss: 0.2079 - accuracy: 0.9219 - val_loss: 0.3876 - val_accuracy: 0.8885
```

```

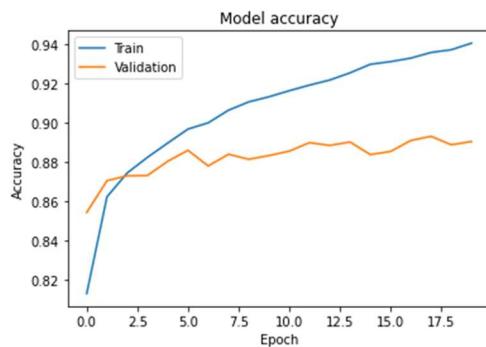
Epoch 14/20
48000/48000 [=====] - 6s 116us/sample - loss: 0.2022 - accuracy: 0.9256 - val_loss: 0.3645 - val_accur
acy: 0.8903
Epoch 15/20
48000/48000 [=====] - 6s 117us/sample - loss: 0.1889 - accuracy: 0.9300 - val_loss: 0.4088 - val_accur
acy: 0.8838
Epoch 16/20
48000/48000 [=====] - 5s 114us/sample - loss: 0.1834 - accuracy: 0.9313 - val_loss: 0.3853 - val_accur
acy: 0.8854
Epoch 17/20
48000/48000 [=====] - 5s 115us/sample - loss: 0.1789 - accuracy: 0.9331 - val_loss: 0.3819 - val_accur
acy: 0.8910
Epoch 18/20
48000/48000 [=====] - 5s 115us/sample - loss: 0.1689 - accuracy: 0.9360 - val_loss: 0.3919 - val_accur
acy: 0.8932
Epoch 19/20
48000/48000 [=====] - 5s 113us/sample - loss: 0.1680 - accuracy: 0.9374 - val_loss: 0.4258 - val_accur
acy: 0.8888
Epoch 20/20
48000/48000 [=====] - 5s 115us/sample - loss: 0.1585 - accuracy: 0.9408 - val_loss: 0.4343 - val_accur
acy: 0.8905
10000/10000 [=====] - 1s 86us/sample - loss: 0.4757 - accuracy: 0.8794
Test accuracy: 0.8794

```

```

In [11]: 1 # Plot training history
2 import matplotlib.pyplot as plt
3
4 plt.plot(history.history['accuracy'])
5 plt.plot(history.history['val_accuracy'])
6 plt.title('Model accuracy')
7 plt.ylabel('Accuracy')
8 plt.xlabel('Epoch')
9 plt.legend(['Train', 'Validation'], loc='upper left')
10 plt.show()

```



LEARNING OUTCOMES:

- Understanding of the MNIST Fashion dataset and its use as a benchmark dataset for image classification tasks.
- Familiarity with the architecture of a neural network for image classification, including the number and size of layers, activation functions, and loss function.
- Ability to preprocess image data for use in a neural network, including normalization and splitting into training and testing sets.
- Ability to implement a neural network using Python and a deep learning library such as TensorFlow or PyTorch for image classification on the MNIST Fashion dataset.